




Estimation

Principle of *Maximum Likelihood Estimation*:
Choose the parameters that maximize the
likelihood of the data

$$\theta^{\text{MLE}} = \arg \max_{\theta} \prod_{t=1}^N P(\mathbf{x}^{(t)} \mid \theta)$$



Which hypothesis is better?



- $\exists$ only a fixed amount of total probability (sums to one)
- MLE pushes as much probability as possible to things we've observed
 - ▶ ...and steals it from things we didn't observe

MLE example

- Generative story: we flip a coin several times; it lands heads w.p. θ , iid over flips
- Data: HTT
- Likelihood:

$$\theta(1-\theta)(1-\theta)$$



- MLE:

$$\ln \theta + \ln((1-\theta)^2) \quad \rightarrow \quad 2 \ln(1-\theta)$$

$$\nabla_{\theta} \ln \theta + \frac{2}{1-\theta}(-1) = 0$$

$$1-\theta = 2\theta$$
$$\theta = 1/3$$

MOO




- Coin flips:
 - ▶ model: iid with probability θ
 - ▶ objective: log likelihood
 - ▶ optimizer: exact (differentiate and set to 0)

Working with probabilities

- Should know how to work with probabilities — recitation / readings / OH / Piazza can help if you're rusty
- Discrete and continuous random variables, events; marginal, joint, and conditional probabilities; expectation and conditional expectation; mean and variance; independence and conditional independence
- Familiarity with common distributions: e.g., Gaussian, Bernoulli, multinomial
- Manipulating probability tables or continuous densities
 - ▶ including conditioning and conditional tables/densities
- Most important tools: factorization; Bayes rule; linearity of expectation
 - ▶ also broadcasting; change of variables (Jacobian rule)

What if we change it a bit?

Very small change
in generative story:
 θ is now uniform
instead of fixed




- Again, a biased coin:
 $p(\text{heads})$ is θ
- Parameter $\theta \in [0,1]$ selected uniformly
 - ▶ the above is what we know before seeing any data (*prior*)
- We now see 5 flips: $X_1, X_2, \dots, X_5$ are HHTHT
- None of the flips depend on any of the others (they are conditionally independent given θ)
- We want the distribution of θ after seeing $X_{1:5}$ (*posterior*)
 - ▶ terminology: prior and posterior are always relative to a given inference — often the posterior from one inference becomes the prior for the next

$$0 \leq \theta \leq 1$$
$$1 \cdot \theta^3 \cdot (1 - \theta)^2$$

***Find
posterior
using Bayes
rule***

- Write out prior $\times$ likelihood, then normalize
 - ▶ $P(\theta \mid X) = P(\theta) P(X \mid \theta) / P(X)$ ← *can skip denominator if we're normalizing anyway*
 - ▶ normalizer measures ***fit***: how well does our entire model class explain our observations

Beta



plot credit: Wikipedia

notice: $\text{Beta}(1, 1)$
is uniform


- Posterior is $P(\theta \mid X_{1:5}) = 60 \theta^3 (1 - \theta)^2$
- This is a **Beta** distribution
 - ▶ in general, $\propto \theta^{\alpha-1} (1 - \theta)^{\beta-1}$, parameters $\alpha, \beta > 0$:
 - ▶ α moves mass toward $\theta = 1$, β moves it toward $\theta = 0$
 - ▶ larger α, β make a sharper distribution

$$\ln 60 + 3 \ln \theta + 2 \ln (1 - \theta)$$

Finding the most probable parameter value

- Posterior = $60 \theta^3 (1 - \theta)^2$
- To find mode, differentiate and set to 0
 - ▶ easier if we take log before derivative
 - ▶ log turns *product* over examples into *sum* over examples, sum rule is simpler than product rule

MAP




- Highest point of posterior (the *mode*) is $\theta = \frac{3}{5}$
 - ▶ if we have to report a single θ , this is a good one
 - ▶ called *maximum a posteriori* or *MAP* estimate
- Can also acknowledge uncertainty: e.g., report the whole distribution, or list of 100 samples from the distribution, or mean and standard deviation, or an interval that contains 95% of its probability

Statistics and estimators


- Very reasonable answer: $\frac{3}{5}$ of our data were heads, so we say $p(H) = \frac{3}{5}$
- Observed fraction of heads is called a *statistic*: a function of our data
 - ▶ other statistics: fraction of tails; $\ln \frac{1 + \text{\#heads}}{1 + \text{\#tails}}$
- A statistic that we use to find θ is called an *estimator*
 - ▶ in this case, $\mathbb{E}(\text{observed fraction}) = \text{true } \theta$
 - ▶ so our estimator is *unbiased*
- Value of estimator is *estimate* — here $\frac{3}{5}$

***Changing
the prior
changes the
estimate***



- Change prior to $p(\theta) = 6\theta(1 - \theta)$ (a Beta distribution)
- Same data: $X_1, X_2, \dots, X_5$ are HHTHT
- Find posterior and MAP for θ after seeing $X_{1:5}$

New MAP



→ $P(\theta | X_{1:5}) = 280 \theta^4 (1 - \theta)^3$ mode is $\theta = \frac{4}{7}$

Influence of prior

- Same data, two priors $\rightarrow$ two different estimates of θ
 - ▶ 2nd prior was narrower, pulls estimate toward center $\frac{1}{2}$
 - ▶ uniform prior in a sense pays more attention to the data (likelihood term), doesn't favor one θ over another

$$\text{Beta}(\alpha, \beta) \propto \theta^{\alpha-1}(1 - \theta)^{\beta-1}$$

Conjugate priors

- Something interesting happened
 - ▶ prior was uniform = Beta(1, 1), posterior was Beta(4, 3)
 - ▶ prior was Beta(2, 2), posterior was Beta(5, 4)
- Not an accident: Beta is the **conjugate prior** for Bernoulli
 - ▶ i.e., when parameter θ has a Beta prior, and when data is Bernoulli(θ), then posterior for θ will also be Beta
- Happens somewhat often:
 - ▶ $x \sim N(\mu, \sigma^2)$: conjugate prior for μ is Normal
 - ▶ $x \sim \text{Exponential}(\lambda)$: conjugate prior for λ is Gamma
- Useful since it simplifies inference:
 - ▶ e.g., observe h Heads and t Tails: add h to α and t to β

***Cf.
maximum
likelihood***

- Compare MAP vs. MLE:
 - ▶ in posterior $p(\theta) p(X_{1:5} | \theta) / p(X_{1:5})$, MAP ignored denominator (can recover by normalizing)
 - ▶ MLE also ignores prior $p(\theta)$, leaving only likelihood
$$\arg \max_{\theta} p(X_{1:5} | \theta)$$
$$= \arg \max_{\theta} p(X_1 | \theta) p(X_2 | \theta) \dots p(X_5 | \theta)$$
$$= \arg \max_{\theta} [\ln p(X_1 | \theta) + \dots + \ln p(X_5 | \theta)]$$
- Like setting prior to uniform, but works even when uniform prior doesn't make sense

***Cf.
maximum
likelihood***

- Both are *consistent*:
 - ▶ under weak assumptions, parameter $\rightarrow$ true value as $| \text{data} | \rightarrow \infty$
- Different advantages:
 - ▶ MLE: no need to think about prior
 - ▶ MAP: a reasonable prior can often make estimates more reliably accurate by reducing variance



MLE for Linear Regression

Optimization Method #2: Closed Form

1. Evaluate

$$\theta^{\text{MLE}} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

2. Return θ^{MLE}



t	θ_1	θ_2	$J(\theta_1, \theta_2)$
MLE	0.59	0.43	0.2

1. You'll work through the view of linear regression as a probabilistic model in the homework!
2. You'll also see how L1 and L2 regularization is equivalent to MAP estimation

Philosophy of probability

- Bayesian view (like the coin w/ Beta prior):
 - ▶ parameters like θ are random variables, just like data
 - ▶ prior and posterior describe our state of knowledge about θ
- Frequentist view (like the coin w/ no prior):
 - ▶ parameters like θ are constants (properties of the world) not random variables
 - ▶ doesn't make sense to ask for a prior or posterior distribution over θ since it's not random
- $\exists$ other philosophies of probability besides frequentist and Bayesian, but these are the two best known

To MAP or not to MAP?

- MAP doesn't make sense for frequentists!
 - ▶ no prior or posterior
- MLE does, and so does regularized MLE
 - ▶ regularized MLE is mathematically equivalent to MAP if regularizer = $\log(\text{prior})$
 - ▶ but justification is different, and MLE can use regularizers that aren't $\log(\text{prior})$
- Frequentist's justifications for either:
 - ▶ consistency, convergence rates, robustness
 - ▶ e.g., find estimator that $\rightarrow \theta$ quickly, even for worst case true distribution of data (minimax convergence rate)

Difference of opinion

- Bayesian critique of frequentist:
 - ▶ you are ignoring useful information (your incoming prior knowledge), so you are leaving performance on the table
- Frequentist critique of Bayesian:
 - ▶ you don't get the best possible minimax convergence rates, so you are leaving performance on the table
- Both are right
 - ▶ optimize different objectives → disagree on best answer
- ML'ers are typically very pragmatic: use Bayesian or frequentist reasoning depending on which leads to a practical computational solution

Learning objectives— MLE, MAP

- How can we define a probabilistic ML model using its generative story?
- What are likelihood and conditional likelihood?
- What is MLE? Why would we expect it to find a good estimate of our model's parameters?
- Find an MLE by differentiating log likelihood
- Work with discrete and continuous distributions, including differences between probability and density, as well as Bayes rule and tools on “working with probabilities” slide
- What is a MAP estimate? Find one by differentiating log prior + log likelihood, or by working with conjugate prior
- How do MLE and MAP compare?
- What are the differences between Bayesian and frequentist reasoning?


Linear regression



10-701 Introduction to Machine Learning
Geoff Gordon

Linear regression

w = slope (in each feature),
 b = intercept



- Regression = supervised prediction of *real* labels
 - ▶ data: features $X^{(t)} \in \mathbb{R}^d$, labels $Y^{(t)} \in \mathbb{R}$, $t = 1 : N$
- *Linear* regression = predict a linear function of features
 - ▶ model: predict $\hat{Y}^{(t)} = w \cdot X^{(t)} = \sum_i w_i X_i^{(t)}$ or $w \cdot X^{(t)} + b$
 - ▶ can omit b if we have a constant feature

Linear regression

$$\hat{y} = w \cdot x + b$$

- Objective:
 - ▶ minimize sum of squared errors:


$$\frac{1}{2} \sum_{t=1}^N \underbrace{(Y^{(t)} - \hat{Y}^{(t)})^2}_{\epsilon^{(t)}}$$

- ▶ plus optional regularizer (see below)

- E.g., minimize

$$L(w) = \frac{1}{2} \sum_{t=1}^N (Y^{(t)} - \hat{Y}^{(t)})^2 + \frac{\lambda}{2} \|w\|_2^2$$

Tiny example



$$L(w, b) = \frac{1}{2} \left[((w \cdot 1 + b) - 5)^2 + (w \cdot 3 + b - 2)^2 + (w \cdot 5 + b - 4)^2 \right]$$


$$\nabla_b L(w, b) = (w \cdot 1 + b) - 5 + (w \cdot 3 + b) - 2 + (w \cdot 5 + b) - 4$$

$$\nabla_w L(w, b) = ((w \cdot 1 + b) - 5) + ((w \cdot 3 + b) - 2) + ((w \cdot 5 + b) - 4) = 0$$

$$9w + 3b = 11$$

$$(w \cdot 1 + b - 5) + (w \cdot 3 + b - 2) + (w \cdot 5 + b - 4) = 0$$

Schematic




$$L(w) = \|Y - Xw\|_2^2 \quad [+ R(w)]$$

features X			labels y
Family History	Resting Blood Pressure	Cholesterol	Cost (k\$)
Yes	Low	Normal	0.1
No	Medium	Normal	0.2
No	Low	Abnormal	10.3
Yes	Medium	Normal	8.7
Yes	High	Abnormal	13.1
No	Low	Normal	0.1
No	Medium	Normal	0.4
Yes	Medium	Abnormal	9.6

- data matrix X , label vector Y ; rows are $X^{(t)}$, $Y^{(t)}$

Schematic




$$L(w) = \|Y - Xw\|_F^2 \quad [+ R(w)]$$

	features X			labels y
	Family History	Resting Blood Pressure	Cholesterol	Cost (k\$)
I	Yes	Low	Normal	0.1
I	No	Medium	Normal	0.2
I	No	Low	Abnormal	10.3
I	Yes	Medium	Normal	8.7
I	Yes	High	Abnormal	13.1
I	No	Low	Normal	0.1
I	No	Medium	Normal	0.4
I	Yes	Medium	Abnormal	9.6

- data matrix X , label vector Y ; rows are $X^{(t)}$, $Y^{(t)}$

Schematic



$$L(w) = \|Y - Xw\|_F^2 \quad [+ R(w)]$$

features X			labels y
Family History	Resting Blood Pressure	Cholesterol	Cost (k\$)
			0.1
			0.2
			10.3
			8.7
			13.1
			0.1
			0.4
			9.6

- data matrix X , label vector Y ; rows are $X^{(t)}$, $Y^{(t)}$

***Turn
categorical
features into
numeric***

Family History	Resting Blood Pressure	Cholesterol	Cost (k\$)
Yes	Low	Normal	0.1
No	Medium	Normal	0.2
No	Low	Abnormal	10.3
Yes	Medium	Normal	8.7
Yes	High	Abnormal	13.1
No	Low	Normal	0.1
No	Medium	Normal	0.4
Yes	Medium	Abnormal	9.6

- Categorical = list of discrete possible values
 - ▶ e.g., weather = {rainy, sunny, windy}

***Turn
categorical
features into
numeric***

Family History	Resting Blood Pressure	Cholesterol	Cost (k\$)
1	Low	Normal	0.1
0	Medium	Normal	0.2
0	Low	Abnormal	10.3
1	Medium	Normal	8.7
1	High	Abnormal	13.1
0	Low	Normal	0.1
0	Medium	Normal	0.4
1	Medium	Abnormal	9.6

- Categorical = list of discrete possible values
 - ▶ e.g., weather = {rainy, sunny, windy}

*Turn
categorical
features into
numeric*

Family History	Resting Blood Pressure	Cholesterol		Cost (k\$)
1	Low	0	Low	0.1
0	Medium	0	Low	0.2
0	Low	1	Normal	10.3
1	Medium	0	Low	8.7
1	High	1	Normal	13.1
0	Low	0	Low	0.1
0	Medium	0	Low	0.4
1	Medium	1	Normal	9.6

- Categorical = list of discrete possible values
 - ▶ e.g., weather = {rainy, sunny, windy}

*Turn
categorical
features into
numeric*

Family History	Resting Blood Pressure	Cholesterol	Cost (k\$)
1	0 0	0	0.1
0	1 0	0	0.2
0	0 0	A 1	10.3
1	1 0	0	8.7
1	0 1	A 1	13.1
0	0 0	0	0.1
0	1 0	0	0.4
1	1 0	A 1	9.6

- Categorical = list of discrete possible values
 - ▶ e.g., weather = {rainy, sunny, windy}

One-hot encoding

*metaphor: like
making one wire
“hot”, i.e.,
having an
applied voltage*

- Tool: *one-hot* encoding
 - ▶ create one binary feature for each value except default
 - ▶ default $\rightarrow (0,0,\dots,0)$
 - ▶ nondefault $\rightarrow$ set one feature to 1, rest to 0
 - ▶ don't have to pick a default (always a 1 in each row)
- E.g.:
 - ▶ no $\rightarrow (0)$ yes $\rightarrow (1)$
 - ▶ rainy $\rightarrow (0,0)$ sunny $\rightarrow (1,0)$ windy $\rightarrow (0,1)$
 - ▶ rainy $\rightarrow (1,0,0)$ sunny $\rightarrow (0,1,0)$ windy $\rightarrow (0,0,1)$

MOO



- Linear regression
 - ▶ model = linear functions of features (+ constant)
 - ▶ objective = sum of squared errors (+ regularizer)
 - ▶ optimizer = exact (differentiate and set to zero, solve with Gaussian elimination)
 - ▶ or several other options

Regularizers

- Common choices:
 - ▶ $\frac{1}{2} \|w\|_2^2$ called *ridge regression*
 - ▶ $\|w\|_1$ called *LASSO*
 - ▶ if we use both, *elastic net*

- E.g., minimize

$$L(w) = \frac{1}{2} \sum_{t=1}^T (Y^{(t)} - \hat{Y}^{(t)})^2 + \frac{\lambda}{2} \|w\|_2^2 + \beta \|w\|_1$$

Why regularize?

- Limit model complexity, trade underfitting vs. overfitting
 - ▶ underfitting is often called ***bias***: the correct answer is not an option, so we pick something that might be far from it
 - ▶ overfitting shows up as ***variance*** of parameters w : more choices → more variability in which one we select
- So you may see *bias-variance tradeoff* to describe effect of model complexity

Overfitting, bias, variance

```
n = 10
k = 7
X = np.random.normal(size=(n, k))
Y = np.random.normal(size=(n, 1))
w, _, _, _ = np.linalg.lstsq(X, Y, rcond=None)
(np.linalg.norm(w), np.mean((Y - X@w)**2))
```

- What is the smallest MSE this code should return if there were no selection bias?
- What is the actual best w ?

Overfitting, bias, variance

```
n = 10
k = 7
X = np.random.normal(size=(n, k))
Y = np.random.normal(size=(n, 1))
w, _, _, _ = np.linalg.lstsq(X, Y, rcond=None)
(np.linalg.norm(w), np.mean((Y - X@w)**2))
```

- What is the smallest MSE this code should return if there were no selection bias?
- What is the actual best w ?

```
# three runs (n=10, k=7):
```

```
(1.729258520013365, 0.09403950328236421)
(1.9706932994777069, 0.07379394277037196)
(3.591128626783351, 0.47238275606461944)
```

Overfitting, bias, variance

```
n = 10
k = 7
X = np.random.normal(size=(n, k))
Y = np.random.normal(size=(n, 1))
w, _, _, _ = np.linalg.lstsq(X, Y, rcond=None)
(np.linalg.norm(w), np.mean((Y - X@w)**2))
```

- What is the smallest MSE this code should return if there were no selection bias?
- What is the actual best w ?

```
# three runs (n=10, k=7):
```

```
(1.729258520013365, 0.09403950328236421)
(1.9706932994777069, 0.07379394277037196)
(3.591128626783351, 0.47238275606461944)
```

```
# three runs (n=20, k=2):
```

```
(0.20006773436081007, 0.8932440732695813)
(0.3765035073938464, 0.9200354120567253)
(0.5764201246357761, 0.7614382010070438)
```

Effect of regularization

- Regularization attempts to limit overfitting by preferring simpler models (smaller weights)

three runs (n=10, k=7, lambda=0):

(1.729258520013365, 0.09403950328236421)
(1.9706932994777069, 0.07379394277037196)
(3.591128626783351, 0.47238275606461944)


three runs (n=10, k=7, lambda=1):

(0.5961635084865388, 0.20898544573161687)
(0.6399685598204747, 0.30440651674916513)
(0.9915710819788339, 0.19630426412988047)

three runs (n=10, k=7, lambda=10):


(0.18652470206125477, 0.20138831819271244)
(0.3682010478765221, 0.4279430122544764)
(0.46978624375583933, 0.44861948640322524)

Model complexity



- Measure complexity of our model class as the number of *effectively distinct* weight vectors
 - ▶ if two weight vectors are very close together, they make almost the same predictions
 - ▶ count only vectors in a grid w/ separation ϵ
- More w s if we have *more features* or *bigger weights*
 - ▶ regularization can limit both


Model complexity



- Measure complexity of our model class as the number of *effectively distinct* weight vectors
 - ▶ if two weight vectors are very close together, they make almost the same predictions
 - ▶ count only vectors in a grid w/ separation ϵ
- More w s if we have *more features* or *bigger weights*
 - ▶ regularization can limit both

Model selection

Earlier we said to limit overfitting by selecting among model classes of different complexity/size — regularization doesn't look quite like this!



- Limit Euclidean norm of weight vector:


$$\min_w L(w) = \frac{1}{2} \sum_{t=1}^N (Y^{(t)} - X^{(t)} \cdot w)^2 \quad \text{s.t.} \quad \|w\|_2 \leq C$$

- Model selection = picking C

Penalty vs. constraint

Varying C or λ gets the same set of models!

Idea: it's easier to pick a good λ ahead of time




- Instead of constrained minimization

$$\min_w L(w) = \frac{1}{2} \sum_{t=1}^N (Y^{(t)} - X^{(t)} \cdot w)^2 \quad \text{s.t.} \quad \|w\|_2 \leq C$$


we can try to directly trade off $L(w)$ and $\|w\|_2$ in optimizer (avoid outer model selection loop):

$$\min_w L(w) = \frac{1}{2} \sum_{t=1}^N (Y^{(t)} - X^{(t)} \cdot w)^2 + \frac{\lambda}{2} \|w\|_2^2$$

***L_1 norm and
count of
nonzeros
are related***



***L_1 norm and
count of
nonzeros
are related***



Pre- processing

- Linear models are extremely useful, but can be sensitive to how the features are expressed
 - ▶ also true of nonlinear models
- Extremely common to modify features before training — called *pre-processing*
 - ▶ e.g., one-hot encoding of categorical features
- Large toolbox of possible ways to pre-process
 - ▶ each tool is designed to fix a potential problem
 - ▶ use any or all of them, as the situation calls for
 - ▶ but don't use them unless they're needed
- Bundle together into a *feature transform*

Ordered features

- Sometimes discrete features have a built-in order
 - ▶ {elementary school, middle school, high school, undergrad, MS, PhD}
- Could use one-hot, but it would erase the order
- Instead, *thermometer code*:
 - ▶ use k features for k values, but turn on the first j of them to encode the j th value
 - ▶ e.g., high school $\rightarrow (1,1,1,0,0,0)$
 - ▶ analogy to the mercury rising in an old-fashioned thermometer
- Can also omit a feature for the first value (assign all zeros) and have $k - 1$ features

Linearly dependent features

- Sometimes a feature will be accidentally duplicated (e.g., two sources provided it w/ different names, or we added a constant feature when one was there already)
- More subtly, a feature might be the sum or difference of two others — e.g., constant vs. one-hot w/o default
- This is no problem w/ regularization
 - ▶ L_2 : if k copies of a feature, we'll assign $\frac{1}{k}$ weight to each
 - ▶ L_1 : any way to split weight is optimal, often prefer sparse
 - ▶ other: will usually do something sane
- But without regularization, optimum might be undefined!
- Diagnose by testing rank of feature matrix X
 - ▶ fix by finding features to drop, or by PCA (later)

Feature range

- Features can have very different typical range of values
 - ▶ e.g., a yes/no question vs. the population of a city
 - ▶ e.g., units: measure distance in mm vs. Mpc
- If no regularization, linear models don't care
 - ▶ multiply feature by 73, regression will divide weight by 73
- But regularization changes this (why?)
- And some optimizers can break (e.g., SGD)
- Two common ways this can happen: location and scale

Location

- Remove location differences by subtracting feature means
 - ▶ feature X_i^t becomes $X_i^t - \hat{\mu}_i$ (called **centering**)
 - ▶ here $\hat{\mu}_i$ is an estimate of the mean (or could use other statistics such as median or trimmed mean)
- Note: must compute $\hat{\mu}_i$ from *training data only*
 - ▶ don't use validation or test data: can't in deployed version, so doing it during development gives wrong answers
 - ▶ $\hat{\mu}_i$ becomes part of our model
- Alternate approach: omit regularization on constant coefficient b in $w \cdot X^t + b$
 - ▶ then b is free to compensate for differences in mean

Scale

- Remove scale differences by dividing by feature scale
 - ▶ feature X_i^t becomes $\frac{1}{\hat{\sigma}_i} X_i^t$ where $\hat{\sigma}_i$ is estimate of scale
 - ▶ e.g., the standard deviation (square root of variance)
 - ▶ or the max absolute value, interquartile range, ...
- As above, must compute $\hat{\sigma}_i$ on *training data only*, and $\hat{\sigma}_i$ becomes part of our model
- If we're transforming location and scale, do location first
 - ▶ scale is more meaningful on centered data

Outliers

- Sometimes a single example will completely dominate the size of a feature
 - ▶ e.g., weight: most examples are mice, one is elephant
 - ▶ or typo: forgot separator, so “030, 001” → “030001”
- Possible fixes:
 - ▶ filter out (delete) the outlying examples
 - ▶ clip the feature (project onto an estimate of the “reasonable” range, e.g., elephant → heaviest mouse)
 - ▶ CDF transform (rank transform): replace feature value by its rank in the training set, then scale to $[0,1]$
 - ▶ whichever we use, has to depend only on training set, becomes part of our model

And more...

- Try to fill in missing feature values
- Encode examples w/ a foundation model
- Text → word part embeddings
- Random projection
- Vector quantization
- Positional encoding
- ...

To process or not to process?


- Tempting to apply every possible preprocessing trick just to be sure
- Unfortunately it's not this simple
 - ▶ can lead to overfitting (parameters: location, scale, CDF)
 - ▶ can erase useful information: maybe the feature that didn't vary much is actually not very informative
 - ▶ interpretability
- Probably best to be conservative:
 - ▶ look for problems that influence the fit: outliers, significant variation in scale, ...
 - ▶ fix at the source if possible
 - ▶ pre-process when necessary

Solving it

- To find best w , several options:
 - ▶ minimize symbolically by setting $\nabla_w L$ to 0
 - ▶ sub-choices: QR, SVD, Cholesky
 - ▶ iterative optimizer
 - ▶ sub-choices: stochastic vs. full gradient
 - ▶ 1st-order like SGD or Adam
 - ▶ 2nd-order like Newton
 - ▶ in-between like L-BFGS
 - ▶ tools like momentum, rescaling
 - ▶ choice of covariance form or kernel form

Exact solution for regression and ridge regression

$$\begin{aligned} X &\in \mathbb{R}^{N \times d} \\ Y &\in \mathbb{R}^N \\ w &\in \mathbb{R}^d \end{aligned}$$



- Predictions are $\hat{Y} = Xw$
- Objective is $\frac{\lambda}{2} \|w\|^2 + \frac{1}{2} \|Y - Xw\|^2$
- Gradient is $\lambda w + X^\top (Xw - Y)$
- Exact solution is $(\lambda I + X^\top X)w = X^\top Y$
 - ▶ implemented in `torch.linalg.lstsq`


*called **normal
equations***

Exact solution for regression and ridge regression

$$X \in \mathbb{R}^{N \times d}$$

$$Y \in \mathbb{R}^N$$

$$w \in \mathbb{R}^d$$



- Prediction

- Objective

- Gradient

- Exact solution

Why use `torch.linalg.lstsq`? Handles (many) special cases: over- or under-determined, rank-deficient, poor conditioning

It would be several days of lectures to describe algorithms and best practices for all of these

*Please remember: don't write `inv()` unless you are **sure** none of the special cases will bite you*

- ▶ implemented in `torch.linalg.lstsq`

called **normal equations**